Package 'clugenr'

August 19, 2023

Title Multidimensional Cluster Generation Using Support Lines

Version 1.0.1

Description An implementation of the clugen algorithm for generating multidimensional clusters with arbitrary distributions. Each cluster is supported by a line segment, the position, orientation and length of which guide where the respective points are placed. This package is described in Fachada & de Andrade (2023) <doi:10.1016/j.knosys.2023.110836>.

Depends R (>= 3.6.0)

Imports mathjaxr

Suggests devtools,
ggplot2,
knitr,
lintr,
patchwork,
rgl,
rmarkdown,
roxygen2,
testthat (>= 3.0.0)

RdMacros mathjaxr

License MIT + file LICENSE

URL https://github.com/clugen/clugenr

BugReports https://github.com/clugen/clugenr/issues

Encoding UTF-8

LazyData true

Roxygen list(markdown = TRUE)

RoxygenNote 7.2.3

Config/testthat/edition 3

VignetteBuilder knitr

1

R topics documented:

---

angle_btw   *Angle between two n-dimensional vectors.*

Description

Typically, the angle between two vectors `v1` and `v2` can be obtained with:

```
acos((v1 %*% v2) / (norm(v1, "2") * norm(v2, "2")))
```

However, this approach is numerically unstable. The version provided here is numerically stable and based on the Angle Between Vectors Julia package by Jeffrey Sarnoff (MIT license), implementing an algorithm provided by Prof. W. Kahan in these notes (see page 15).

Usage

```
angle_btw(v1, v2)
```

Arguments

v1	First vector.
v2	Second vector.

Value

Angle between `v1` and `v2` in radians.

Examples

```
angle_btw(c(1.0, 1.0, 1.0, 1.0), c(1.0, 0.0, 0.0, 0.0)) * 180 / pi
```

---

angle_deltas	*Get angles between average cluster direction and cluster-supporting lines*

Description

Determine the angles between the average cluster direction and the cluster-supporting lines. These angles are obtained from a wrapped normal distribution ($\mu = 0$, $\sigma = $ `angle_disp`) with support in the interval $[-\pi/2, \pi/2]$. Note this is different from the standard wrapped normal distribution, the support of which is given by the interval $[-\pi, \pi]$.

Usage

```
angle_deltas(num_clusters, angle_disp)
```

Arguments

num_clusters   Number of clusters.

angle_disp   Angle dispersion, in radians.

Value

Angles between the average cluster direction and the cluster-supporting lines, given in radians in the interval $[-\pi/2, \pi/2]$

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
set.seed(123)
arad <- angle_deltas(4, pi / 8) # Angle dispersion of 22.5 degrees
arad                            # What angles deltas did we get?
arad * 180 / pi                 #  Show angle deltas in degrees
```

---

clucenters	*Determine cluster centers using the uniform distribution*

Description

Determine cluster centers using the uniform distribution, taking into account the number of clusters (num_clusters) and the average cluster separation (clu_sep).

More specifically, let c =num_clusters, $\mathbf{s}$ =clu_sep, $\mathbf{o}$ =clu_offset, n =length(clu_sep) (i.e., number of dimensions). Cluster centers are obtained according to the following equation:

$$\mathbf{C} = c\mathbf{U} \cdot \mathrm{diag}(\mathbf{s}) + \mathbf{1}\,\mathbf{o}^T$$

where $\mathbf{C}$ is the $c \times n$ matrix of cluster centers, $\mathbf{U}$ is an $c \times n$ matrix of random values drawn from the uniform distribution between -0.5 and 0.5, and $\mathbf{1}$ is an $c \times 1$ vector with all entries equal to 1.

Usage

```
clucenters(num_clusters, clu_sep, clu_offset)
```

Arguments

num_clusters	Number of clusters.
clu_sep	Average cluster separation ($n \times 1$ vector).
clu_offset	Cluster offsets ($n \times 1$ vector).

Value

A $c \times n$ matrix containing the cluster centers.

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
set.seed(321)
clucenters(3, c(30, 10), c(-50,50))
```

---

clugen   *Generate multidimensional clusters*

Description

This is the main function of **clugenr**, and possibly the only function most users will need.

Usage

```
clugen(
  num_dims,
  num_clusters,
  num_points,
  direction,
  angle_disp,
  cluster_sep,
  llength,
  llength_disp,
  lateral_disp,
  allow_empty = FALSE,
  cluster_offset = NA,
  proj_dist_fn = "norm",
  point_dist_fn = "n-1",
  clusizes_fn = clusizes,
  clucenters_fn = clucenters,
  llengths_fn = llengths,
  angle_deltas_fn = angle_deltas,
  seed = NA
)
```

Arguments

num_dims	Number of dimensions.
num_clusters	Number of clusters to generate.
num_points	Total number of points to generate.
direction	Average direction of the cluster-supporting lines. Can be a vector of length num_dims (same direction for all clusters) or a matrix of size num_clusters x num_dims (one direction per cluster).
angle_disp	Angle dispersion of cluster-supporting lines (radians).
cluster_sep	Average cluster separation in each dimension (vector of length num_dims).
llength	Average length of cluster-supporting lines.
llength_disp	Length dispersion of cluster-supporting lines.
lateral_disp	Cluster lateral dispersion, i.e., dispersion of points from their projection on the cluster-supporting line.
allow_empty	Allow empty clusters? FALSE by default.
cluster_offset	Offset to add to all cluster centers (vector of length num_dims). By default there will be no offset.
proj_dist_fn	Distribution of point projections along cluster-supporting lines, with three possible values:

- "norm" (default): Distribute point projections along lines using a normal distribution ($\mu = line\_center$, $\sigma = $ llength/6).
- "unif": Distribute points uniformly along the line.

- User-defined function, which accepts two parameters, line length (double) and number of points (integer), and returns a vector containing the distance of each point projection to the center of the line. For example, the "norm" option roughly corresponds to function(l, n) stats::rnorm(n, sd = l / 6).

point_dist_fn   Controls how the final points are created from their projections on the cluster-supporting lines, with three possible values:

- "n-1" (default): Final points are placed on a hyperplane orthogonal to the cluster-supporting line, centered at each point's projection, using the normal distribution ($\mu = 0$, $\sigma =$ lateral_disp). This is done by the clupoints_n_1 function.
- "n": Final points are placed around their projection on the cluster-supporting line using the normal distribution ($\mu = 0$, $\sigma =$ lateral_disp). This is done by the clupoints_n function.
- User-defined function: The user can specify a custom point placement strategy by passing a function with the same signature as clupoints_n_1 and clupoints_n.

clusizes_fn   Distribution of cluster sizes. By default, cluster sizes are determined by the clusizes function, which uses the normal distribution ($\mu =$ num_points/num_clusters, $\sigma = \mu/3$), and assures that the final cluster sizes add up to num_points. This parameter allows the user to specify a custom function for this purpose, which must follow clusizes signature. Note that custom functions are not required to strictly obey the num_points parameter. Alternatively, the user can specify a vector of cluster sizes directly.

clucenters_fn   Distribution of cluster centers. By default, cluster centers are determined by the clucenters function, which uses the uniform distribution, and takes into account the num_clusters and cluster_sep parameters for generating well-distributed cluster centers. This parameter allows the user to specify a custom function for this purpose, which must follow clucenters signature. Alternatively, the user can specify a matrix of size num_clusters x num_dims with the exact cluster centers.

llengths_fn   Distribution of line lengths. By default, the lengths of cluster-supporting lines are determined by the llengths function, which uses the folded normal distribution ($\mu =$ llength, $\sigma =$ llength_disp). This parameter allows the user to specify a custom function for this purpose, which must follow llengths signature. Alternatively, the user can specify a vector of line lengths directly.

angle_deltas_fn

Distribution of line angle differences with respect to direction. By default, the angles between the main direction of each cluster and the final directions of their cluster-supporting lines are determined by the angle_deltas function, which uses the wrapped normal distribution ($\mu = 0$, $\sigma =$ angle_disp) with support in the interval $[-\pi/2, \pi/2]$. This parameter allows the user to specify a custom function for this purpose, which must follow angle_deltas signature. Alternatively, the user can specify a vector of angle deltas directly.

seed   An integer used to initialize the PRNG, allowing for reproducible results. If specified, seed is simply passed to set.seed.

Details

If a custom function was given in the `clusizes_fn` parameter, it is possible that `num_points` may have a different value than what was specified in the `num_points` parameter.

The terms "average" and "dispersion" refer to measures of central tendency and statistical dispersion, respectively. Their exact meaning depends on the optional arguments.

Value

A named list with the following elements:

- `points`: A `num_points` x `num_dims` matrix with the generated points for all clusters.
- `clusters`: A `num_points` factor vector indicating which cluster each point in `points` belongs to.
- `projections`: A `num_points` x `num_dims` matrix with the point projections on the cluster-supporting lines.
- `sizes`: A `num_clusters` x 1 vector with the number of points in each cluster.
- `centers`: A `num_clusters` x `num_dims` matrix with the coordinates of the cluster centers.
- `directions`: A `num_clusters` x `num_dims` matrix with the final direction of each cluster-supporting line.
- `angles`: A `num_clusters` x 1 vector with the angles between the cluster-supporting lines and the main direction.
- `lengths`: A `num_clusters` x 1 vector with the lengths of the cluster-supporting lines.

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
# 2D example
x <- clugen(2, 5, 1000, c(1, 3), 0.5, c(10, 10), 8, 1.5, 2)
graphics::plot(x$points, col = x$clusters, xlab = "x", ylab = "y", asp = 1)
# 3D example
x <- clugen(3, 5, 1000, c(2, 3, 4), 0.5, c(15, 13, 14), 7, 1, 2)
```

---

clumerge   *Merges the fields (specified in* fields*) of two or more data sets*

Description

Merges the fields (specified in `fields`) of two or more data sets (passed as lists). The fields to be merged need to have the same number of columns. The corresponding merged field will contain the rows of the fields to be merged, and will have a common "supertype".

Usage

```
clumerge(..., fields = c("points", "clusters"), clusters_field = "clusters")
```

Arguments

`...`	One or more cluster data sets (in the form of lists) whose `fields` are to be merged.
`fields`	Fields to be merged, which must exist in the data sets given in `...`.
`clusters_field`	Field containing the integer cluster labels. If specified, cluster assignments in individual datasets will be updated in the merged dataset so that clusters are considered separate.

Details

The `clusters_field` parameter specifies a field containing integers that identify the cluster to which the respective points belongs to. If `clusters_field` is specified (by default it's specified as `"clusters"`), cluster assignments in individual datasets will be updated in the merged dataset so that clusters are considered separate. This parameter can be set to `NA`, in which case no field will be considered as a special cluster assignments field.

This function can be used to merge data sets generated with the [clugen](#) function, by default merging the `points` and `clusters` fields in those data sets. It also works with arbitrary data by specifying alternative fields in the `fields` parameter. It can be used, for example, to merge third-party data with [clugen](#)-generated data.

Value

A list whose fields consist of the merged fields in the original data sets.

Examples

```
a <- clugen(2, 5, 100, c(1, 3), 0.5, c(10, 10), 8, 1.5, 2)
b <- clugen(2, 3, 250, c(-1, 3), 0.5, c(13, 14), 7, 1, 2)
ab <- clumerge(a, b)
```

---

clupoints_n	*Create points from their projections on a cluster-supporting line*

Description

Each point is placed around its projection using the normal distribution ($\mu = 0$, $\sigma = $ `lat_disp`).

Usage

```
clupoints_n(projs, lat_disp, line_len, clu_dir, clu_ctr)
```

Arguments

projs	Point projections on the cluster-supporting line ($p \times n$ matrix).
lat_disp	Standard deviation for the normal distribution, i.e., cluster lateral dispersion.
line_len	Length of cluster-supporting line (ignored).
clu_dir	Direction of the cluster-supporting line.
clu_ctr	Center position of the cluster-supporting line (ignored).

Details

This function's main intended use is by the main clugen function, generating the final points when the point_dist_fn parameter is set to "n".

Value

Generated points ($p \times n$ matrix).

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
set.seed(123)
ctr <- c(0, 0)
dir <- c(1, 0)
pdist <- c(-0.5, -0.2, 0.1, 0.3)
proj <- points_on_line(ctr, dir, pdist)
clupoints_n(proj, 0.01, NA, dir, NA)
```

---

clupoints_n_1 *Create points from their projections on a cluster-supporting line*

Description

Each point is placed on a hyperplane orthogonal to that line and centered at the point's projection, using the normal distribution ($\mu = 0, \sigma = $ lat_disp).

Usage

```
clupoints_n_1(projs, lat_disp, line_len, clu_dir, clu_ctr)
```

Arguments

projs	Point projections on the cluster-supporting line ($p \times n$ matrix).
lat_disp	Standard deviation for the normal distribution, i.e., cluster lateral dispersion.
line_len	Length of cluster-supporting line (ignored).
clu_dir	Direction of the cluster-supporting line.
clu_ctr	Center position of the cluster-supporting line (ignored).

Details

This function's main intended use is by the main [clugen](#) function, generating the final points when
the `point_dist_fn` parameter is set to `"n-1"`.

Value

Generated points ($p \times n$ matrix).

Note

This function is stochastic. For reproducibility set a PRNG seed with [set.seed](#).

Examples

```
set.seed(123)
ctr <- c(0, 0)
dir <- c(1, 0)
pdist <- c(-0.5, -0.2, 0.1, 0.3)
proj <- points_on_line(ctr, dir, pdist)
clupoints_n_1(proj, 0.1, NA, dir, NA)
```

---

clupoints_n_1_template

Create points from their projections on a cluster-supporting line

---

Description

Generate points from their n-dimensional projections on a cluster-supporting line, placing each
point on a hyperplane orthogonal to that line and centered at the point's projection. The function
specified in `dist_fn` is used to perform the actual placement.

Usage

```
clupoints_n_1_template(projs, lat_disp, clu_dir, dist_fn)
```

Arguments

`projs`	Point projections on the cluster-supporting line ($p \times n$ matrix).
`lat_disp`	Dispersion of points from their projection.
`clu_dir`	Direction of the cluster-supporting line (unit vector).
`dist_fn`	Function to place points on a second line, orthogonal to the first.

Details

This function is used internally by [clupoints_n_1](#) and may be useful for constructing user-defined
final point placement strategies for the `point_dist_fn` parameter of the main [clugen](#) function.

Value

Generated points ($p \times n$ matrix).

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
set.seed(123)
ctr <- c(0, 0)
dir <- c(1, 0)
pdist <- c(-0.5, -0.2, 0.1, 0.3)
proj <- points_on_line(ctr, dir, pdist)
clupoints_n_1_template(proj, 0, dir, function(p, l) stats::runif(p))
```

---

clusizes   *Determine cluster sizes, i.e., the number of points in each cluster*

Description

Cluster sizes are determined using the normal distribution (μ = num_points / num_clusters, $\sigma = \mu/3$), and then assuring that the final cluster sizes add up to num_points via the fix_num_points function.

Usage

```
clusizes(num_clusters, num_points, allow_empty)
```

Arguments

num_clusters	Number of clusters.
num_points	Total number of points.
allow_empty	Allow empty clusters?

Value

Number of points in each cluster (vector of length num_clusters).

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
set.seed(123)
sizes <- clusizes(4, 1000, TRUE)
sizes
sum(sizes)
```

---

fix_empty   *Certify that, given enough points, no clusters are left empty*

Description

Certifies that, given enough points, no clusters are left empty. This is done by removing a point from the largest cluster and adding it to an empty cluster while there are empty clusters. If the total number of points is smaller than the number of clusters (or if the allow_empty parameter is set to TRUE), this function does nothing.

Usage

```
fix_empty(clu_num_points, allow_empty = FALSE)
```

Arguments

clu_num_points   Number of points in each cluster (vector of size c), where c is the number of clusters.

allow_empty   Allow empty clusters?

Details

This function is used internally by clusizes and might be useful for custom cluster sizing implementations given as the clusizes_fn parameter of the main clugen function.

Value

Number of points in each cluster, after being fixed by this function (vector of size c).

Examples

```
clusters <- c(3, 4, 5, 0, 0)    # A vector with some empty elements
clusters <- fix_empty(clusters) # Apply this function
clusters                        # Check that there's no more empty elements
```

---

fix_num_points   *Certify that array values add up to a specific total*

Description

Certifies that the values in the clu_num_points array, i.e. the number of points in each cluster, add up to num_points. If this is not the case, the clu_num_points array is modified in-place, incrementing the value corresponding to the smallest cluster while sum(clu_num_points) < num_points, or decrementing the value corresponding to the largest cluster while sum(clu_num_points) > num_points.

Usage

```
fix_num_points(clu_num_points, num_points)
```

Arguments

clu_num_points   Number of points in each cluster (vector of size c), where c is the number of clusters.

num_points   The expected total number of points.

Details

This function is used internally by [clusizes](#) and might be useful for custom cluster sizing implementations given as the clusizes_fn parameter of the main [clugen](#) function.

Value

Number of points in each cluster, after being fixed by this function.

Examples

```
clusters <- c(1, 6, 3)                  # 10 total points
clusters <- fix_num_points(clusters, 12) # But we want 12 total points
clusters                                 # Check that we now have 12 points
```

---

llengths   *Determine length of cluster-supporting lines*

Description

Line lengths are determined using the folded normal distribution ($\mu = $ llength, $\sigma = $ llength_disp).

Usage

```
llengths(num_clusters, llength, llength_disp)
```

Arguments

num_clusters   Number of clusters.

llength   Average line length.

llength_disp   Line length dispersion.

Value

Lengths of cluster-supporting lines (vector of size num_clusters).

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
set.seed(123)
llengths(4, 20, 3.5)
```

---

points_on_line   *Determine coordinates of points on a line*

Description

Determine coordinates of points on a line with center and direction, based on the distances from the center given in dist_center.

This works by using the vector formulation of the line equation assuming direction is a n-dimensional unit vector. In other words, considering $\mathbf{d} =$ as.matrix(direction) ($n \times 1$ vector), $\mathbf{c} =$ as.matrix(center) ($n \times 1$ vector), and $\mathbf{w} =$ as.matrix(dist_center) ($p \times 1$ vector), the coordinates of points on the line are given by:

$$\mathbf{P} = \mathbf{1}\,\mathbf{c}^T + \mathbf{w}\mathbf{d}^T$$

where $\mathbf{P}$ is the $p \times n$ matrix of point coordinates on the line, and $\mathbf{1}$ is a $p \times 1$ vector with all entries equal to 1.

Usage

```
points_on_line(center, direction, dist_center)
```

Arguments

center	Center of the line (n-component vector).
direction	Line direction (n-component unit vector).
dist_center	Distance of each point to the center of the line (n-component vector, where n is the number of points).

Value

Coordinates of points on the specified line ($p \times n$ matrix).

Examples

```
points_on_line(c(5, 5), c(1, 0), seq(-4, 4, length.out=5)) # 2D, 5 points

points_on_line(c(-2, 0, 0, 2), c(0, 0, -1, 0), c(10, -10)) # 4D, 2 points
```

---

rand_ortho_vector *Get a random unit vector orthogonal to* u*.*

Description

Get a random unit vector orthogonal to u.

Usage

```
rand_ortho_vector(u)
```

Arguments

u   A unit vector.

Value

A random unit vector orthogonal to u.

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
r <- stats::runif(3)     # Get a random 3D vector
r <- r / norm(r, "2")    # Normalize it
o <- rand_ortho_vector(r) # Get a random unit vector orthogonal to r
r %*% o             # Check that r and o are orthogonal (result should be ~0)
```

---

rand_unit_vector *Get a random unit vector with* num_dims *components.*

Description

Get a random unit vector with num_dims components.

Usage

```
rand_unit_vector(num_dims)
```

Arguments

num_dims   Number of components in vector (i.e. vector size).

Value

A random unit vector with num_dims components.

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
r <- rand_unit_vector(4)
norm(r, "2")
```

---

rand_vector_at_angle   *Get a random unit vector at a given angle with another vector.*

Description

Get a random unit vector which is at angle radians of vector u. Note that u is expected to be a unit vector itself.

Usage

```
rand_vector_at_angle(u, angle)
```

Arguments

u	Unit vector with n components.
angle	Angle in radians.

Value

Random unit vector with n components which is at angle radians with vector u.

Note

This function is stochastic. For reproducibility set a PRNG seed with set.seed.

Examples

```
u <- c(1.0, 0, 0.5, -0.5)          # Define a 4D vector
u <- u / norm(u, "2")              # Normalize the vector
v <- rand_vector_at_angle(u, pi / 4) # Get a vector at 45 degrees
arad <- acos((u %*% v) / norm(u,"2") * norm(v, "2")) # Get angle in radians
arad * 180 / pi # Convert to degrees, should be close to 45 degrees
```

Index